


# Introduction to K8S workload

莊家雋

# Wordpress + MySQL的例子



# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume

# Kubernetes - Pod


- Pod為 Kubernetes 中最小的部署單位 Pod
  - 由一至多個容器所組成
  - 通常將耦合性較高之container放在同一個pod
  - 同一個pod的container會在同一個node
- 共享同樣的 network space
  - Pod內的container可以用localhost溝通
- 共享儲存空間 (volumes)
  - 在Pod定義volume, container可以掛載這些定義好的volume

# Wordpress + MySQL: 一個pod有多個Conatiner

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
  - name: mysql
    image: mysql:5.6
    env:
    - name: MYSQL_ROOT_PASSWORD
      value: Password1234
  - name: wordpress
    image: wordpress:4.8-apache
    env:
    - name: WORDPRESS_DB_HOST
      value: 127.0.0.1
    - name: WORDPRESS_DB_PASSWORD
      value: Password1234
```

# 透過port-forward “暫時” 可以存取


```
$ kubectl port-forward mysql-wordpress-pod 8080:80
```



# Wordpress + MySQL:

## 一個pod只有一個container

需手動找出Pod IP



# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume

# Label & Annotations

```
"environment" : "dev", "environment" : "qa", "environment" : "production"  
  
"tier" : "frontend", "tier" : "backend", "tier" : "cache"  
  
"partition" : "customerA", "partition" : "customerB"
```

- 相同點

- key/value pair
- 每個物件可以同時擁有許多個labels (or annotation)
- 皆位於 object 的 metadata 欄位內

- 不同點

- Label 可以透過selector進行篩選
  - `kubectl get node -l <key>=<value>`
- annotation做為沒有識別用途的標籤
  - 程式讀取annotation的值進行特定工作
  - Eg: `prometheus.io/path: "/metrics"`

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: mysql-pod  
  labels:  
    app: blog  
    tier: mysql  
  annotations:  
    compnay: 'NCHC'  
spec:  
  containers:  
    - name: mysql  
      image: mysql:5.6
```



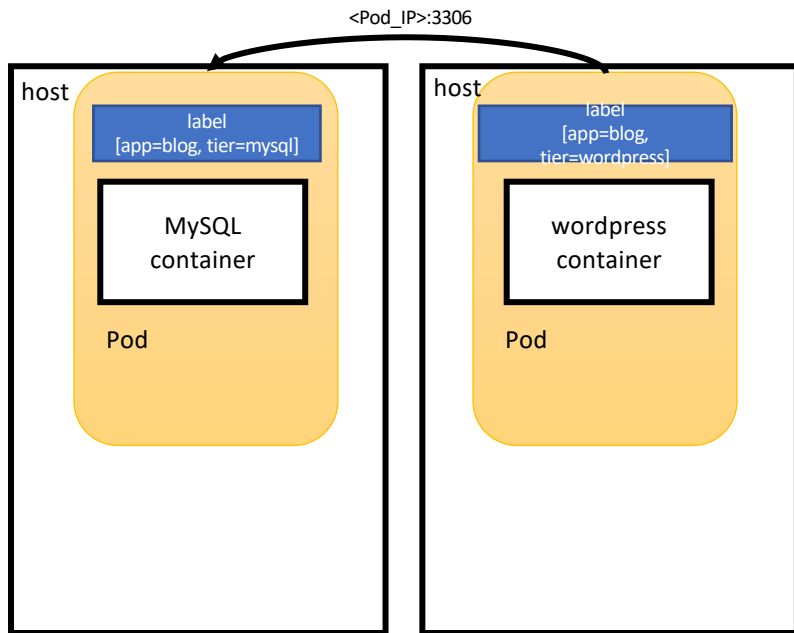
# Label的應用：讓Pod運行在特定節點上

- 在Pod Definition裡使用NodeSelector
- 找出目前node的label
  - `kubectl get node - show-labels`
- 在node上打上label
  - `kubectl label node <name> <key>=<value>`
- 在node上移除label
  - `kubectl label node <name> <key>-`

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
  nodeSelector:
    hardware: high-memory
```

# Wordpress + MySQL: 每個Pod都不同的Label

需手動找出Pod IP



```
$ kubectl get pod -o wide --show-labels  
$ kubectl get pod -o wide -l app=blog
```

# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume

# Kubernetes built-in workload

- Pod 是K8S裡最小的部built-in workload署單位，但是不建議直接使用Pod，而是用built-in workload
  - Deployment
  - Job/CronJob
  - StatefulSet
  - DemonSet
- Built-in workload有相對應的controller協助管理

# Deployments


- Deployment 的功用
  - 確保 Pod 數量(replicas)滿足所設定的值
  - 變動Pod的數量(Scale)
  - 滾動升級(Rolling update)
  - 回滾(Roll back)的機制

Selector

Pod Definition

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql-deploy
  labels:
    app: blog
spec:
  selector:
    matchLabels:
      app: blog
      tier: mysql
  replicas: 1
  template:
    metadata:
      name: mysql-pod
      labels:
        app: blog
        tier: mysql
    spec:
      containers:
        - name: mysql
          image: mysql:5.6
          env:
            - name: MYSQL_ROOT_PASSWORD
              value: Password1234
          volumeMounts:
            - name: mysql-storage
              mountPath: /var/lib/mysql
      volumes:
        - name: mysql-storage
          hostPath:
            path: /tmp/data
```

# Wordpress + MySQL的例子




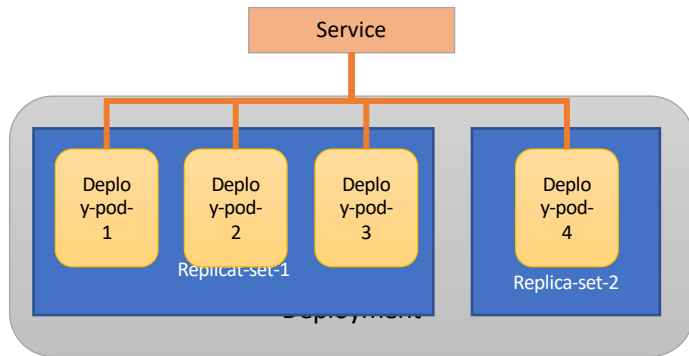


# Deployment 的主要功能

- Scaling a Deployment
  - 擴展pod的數量
- Rolling update (zero down time update)
  - 可以達到無停機服務遷移
- Roll back
  - 當update 完後發現此次的升級造成服務發生不穩定的狀況，可以利用rollback 來回復到先前的狀態。



# Rolling-update



# 常用 Deployment 指令

| Deployment相關指令   | 指令功能                          |   |
|--|-------------------------------|---|
| kubectl <b>scale</b> deploy <deployment> --replicas=n                | Scale deployment 物件的pod數      | } |
| kubectl <b>edit</b> deploy <deployment>                              | 編輯特定deployment物件              |   |
| kubectl <b>set image</b> deployment <deployment> <container>=<image> | 將deployment管理的pod升級到特定image版本 | } |
| kubectl <b>rollout status</b> deploy <deployment>                    | 查詢目前某deployment升級狀況           |   |
| kubectl <b>rollout history</b> deploy <deployment>                   | 查詢目前某deployment升級的歷史紀錄        | } |
| kubectl <b>rollout undo</b> deploy <deployment> --to-revision=n      | 回滾Pod到某個特定版本                  |   |
| kubectl <b>rollout undo</b> deploy <deployment>                      | 回滾Pod到先前一個版本                  |   |

scale

Rolling  
update

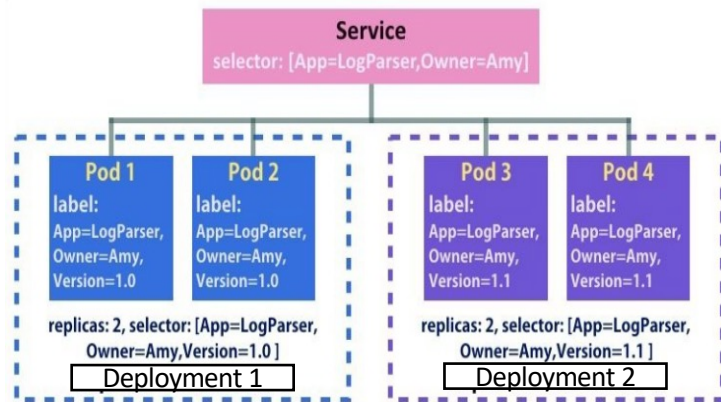

Roll  
back

# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume

# 為什麼需要Service

- Pod的 IP 會變動，不應該透過Pod IP進行存取
- Deployment 有多個replica, 都有各自的Pod IP
- 透過Service的cluster IP，service會將至cluster IP的流量導入不同的pod。
  - 透過Label與selector決定service與pod的對應
  - 可以透過<svc>.<namespace> 的DNS來存取，而不透過IP



# Service

- Service有三種類型
  - ClusterIP、NodePort、LoadBalancer
- Kubernetes 內有三種不同的IP
  - Pod 有 Pod IP
  - Service 有
    - ClusterIP: 建立svc 就會產生
    - External IP: 需由cloud provider動態分配

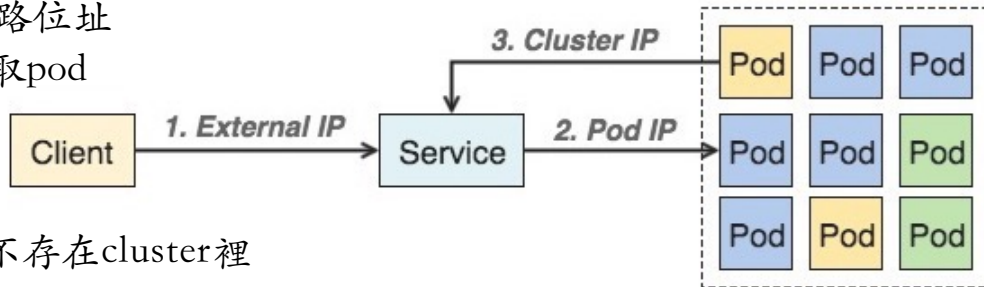
# Kubernetes 的三種IP

- Pod IP

- 每個 Pod 在overlay network上的網路位址
- Cluster 內的 Pod/node 可用此ip存取pod

- Cluster IP

- service在叢集內部的網路位址
- 是一個virtual ip, 不是real ip, 且不存在cluster裡
- 只有Cluster 內的 Pod 可用此ip存取該服務




- External IP

- 透過cloud provider提供
- Service暴露在外的位址, 供外部使用者連線

```
root@ubuntu:~# kubectl get pod -o wide
NAME                                READY    STATUS    RESTARTS   AGE    IP
nfs-provisioner-5bfd94c4b-86g7w    1/1      Running   0           11d    172.20.0.170
root@ubuntu:~# kubectl get svc -n=ingress-nginx
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)
default-http-backend ClusterIP    10.68.90.254  <none>         80/TCP
```


# Service type

- ClusterIP:  $\text{svc.Port} \rightarrow \text{svc.targetPort}$ 
  1. Expose 一個cluster內部可存取的IP
- NodePort:  $\text{nodePort} \rightarrow \text{svc.Port} \rightarrow \text{svc.targetPort}$ 
  1. Expose 一個cluster內部可存取的IP
  2. 在每個node的ip都expose相同的port
- Loadbalancer:  $\text{LB} \rightarrow \text{nodePort} \rightarrow \text{svc.port} \rightarrow \text{svc.targetPort}$ 
  1. Expose 一個cluster內部可存取的IP
  2. 在每個node的ip都expose相同的port
  3. Cloud provider提供一個LB的IP，此IP綁定到expose 的nodePort



# Node Port

- Expose service to outside cluster
- One service uses one nodePort




```
root@ubuntu:~# kubectl get svc -n=ingress-nginx
```

| NAME                 | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)                    | AGE |
|----------------------|-----------|---------------|-------------|----------------------------|-----|
| default-http-backend | ClusterIP | 10.68.90.254  | <none>      | 80/TCP                     | 11d |
| ingress-nginx        | NodePort  | 10.68.163.140 | <none>      | 80:32447/TCP,443:37619/TCP | 10d |



# LoadBalancer

- Provided by cloud provider
- 若沒有Cloud provider, 則會一直pending



| NAME         | TYPE         | CLUSTER-IP    | EXTERNAL-IP | PORT(S)                     | AGE |
|--------------|--------------|---------------|-------------|-----------------------------|-----|
| hub          | ClusterIP    | 10.68.169.105 | <none>      | 8081/TCP                    | 12m |
| proxy-api    | ClusterIP    | 10.68.242.208 | <none>      | 8001/TCP                    | 12m |
| proxy-http   | ClusterIP    | 10.68.144.45  | <none>      | 8000/TCP                    | 12m |
| proxy-public | LoadBalancer | 10.68.89.253  | <pending>   | 80:24837/TCP, 443:22383/TCP | 12m |

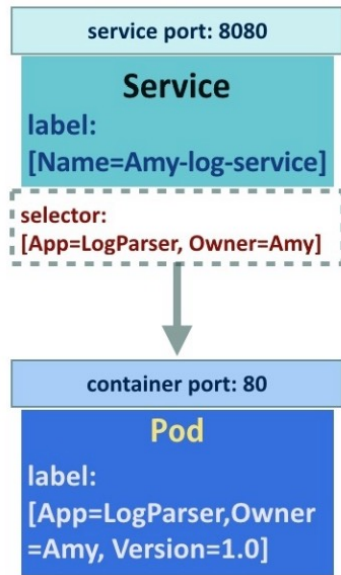
- 用expose建立Service
  - 可以用在pod、deploy

```
$ kubectl expose pod <POD_NAME> --labels="Name=Amy-log-service"
--selector="App=LogParser,Owner=Amy" --port=8080 --target-port=80
--name="my-service"
```

```
$ kubectl expose deploy nginx-deploy --name="service-deploy" --port 80
```

- 用YAML檔描述Service

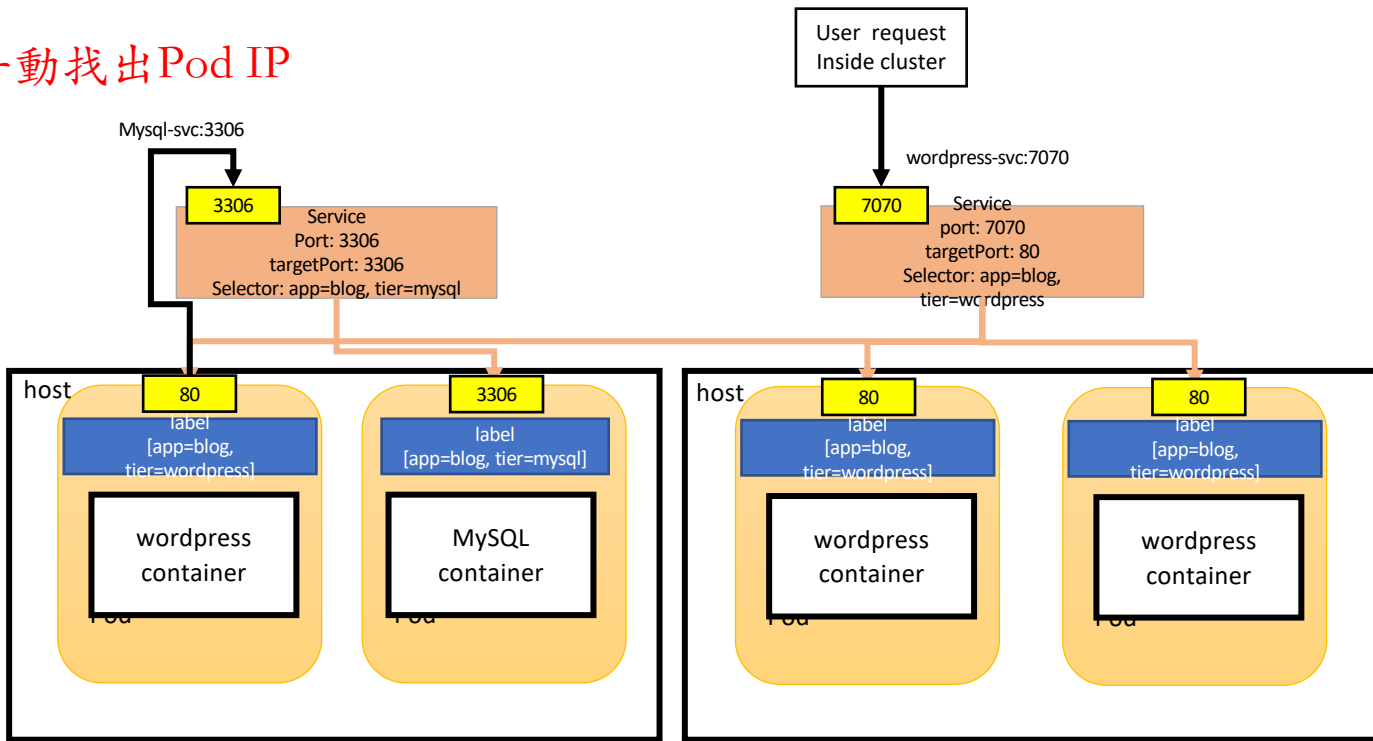
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  label:
    Name: Amy-log-service
selector:
  App: LogParser
  Owner: Amy
spec:
  type: ClusterIP
  ports:
  - port: 8080
    targetPort: 80
```



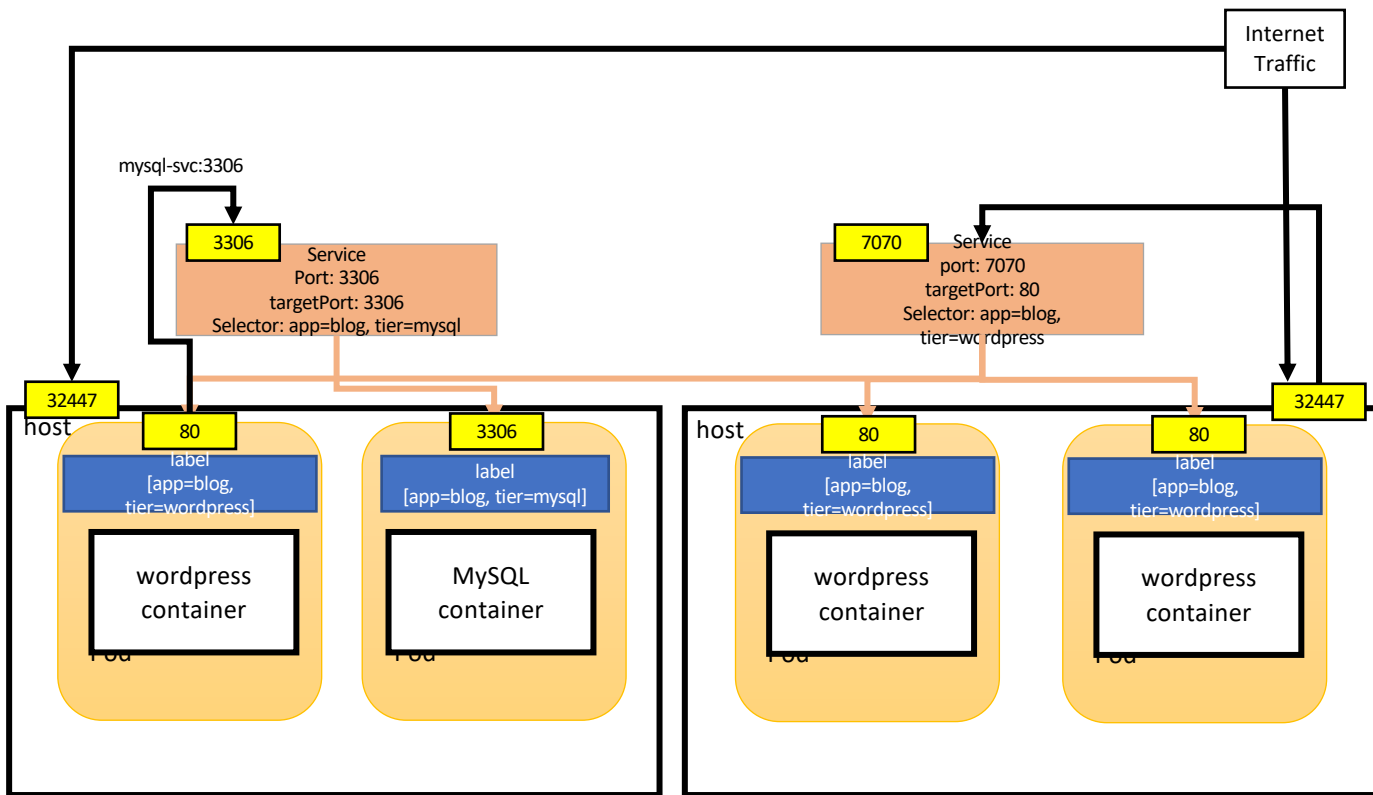
# Wordpress + MySQL:

## 建立ClusterIP的Service, 讓deployment互連

不需手動找出Pod IP



# Wordpress + MySQL的： 建立NodePort的Service讓外部存取



# 如何讓外部存取Pod

- 透過 Service 提供的機制
  - NodePort
  - Load Balancer
- 透過某台主機的網路
  - HostPort
  - 將容器的port和所在節點的port做對應，透過節點ip:port來訪問容器
    - `kubectl apply -f example/04-service/hostnetwork/hostport.yaml`
    - 透過hostIP:8088 存取
- Ingress (後續補充)
  - 只能用在網頁類型服務

# Summary

- Service有ClusterIP, NodePort, LoadBalancer三種
- 透過Service, 可以用<SVC>.<NAMESPACE>存取Pod
- 要讓外部存取pod, 可以使用NodePort, LoadBalancer, Ingress, HostPort

# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume

# Secret

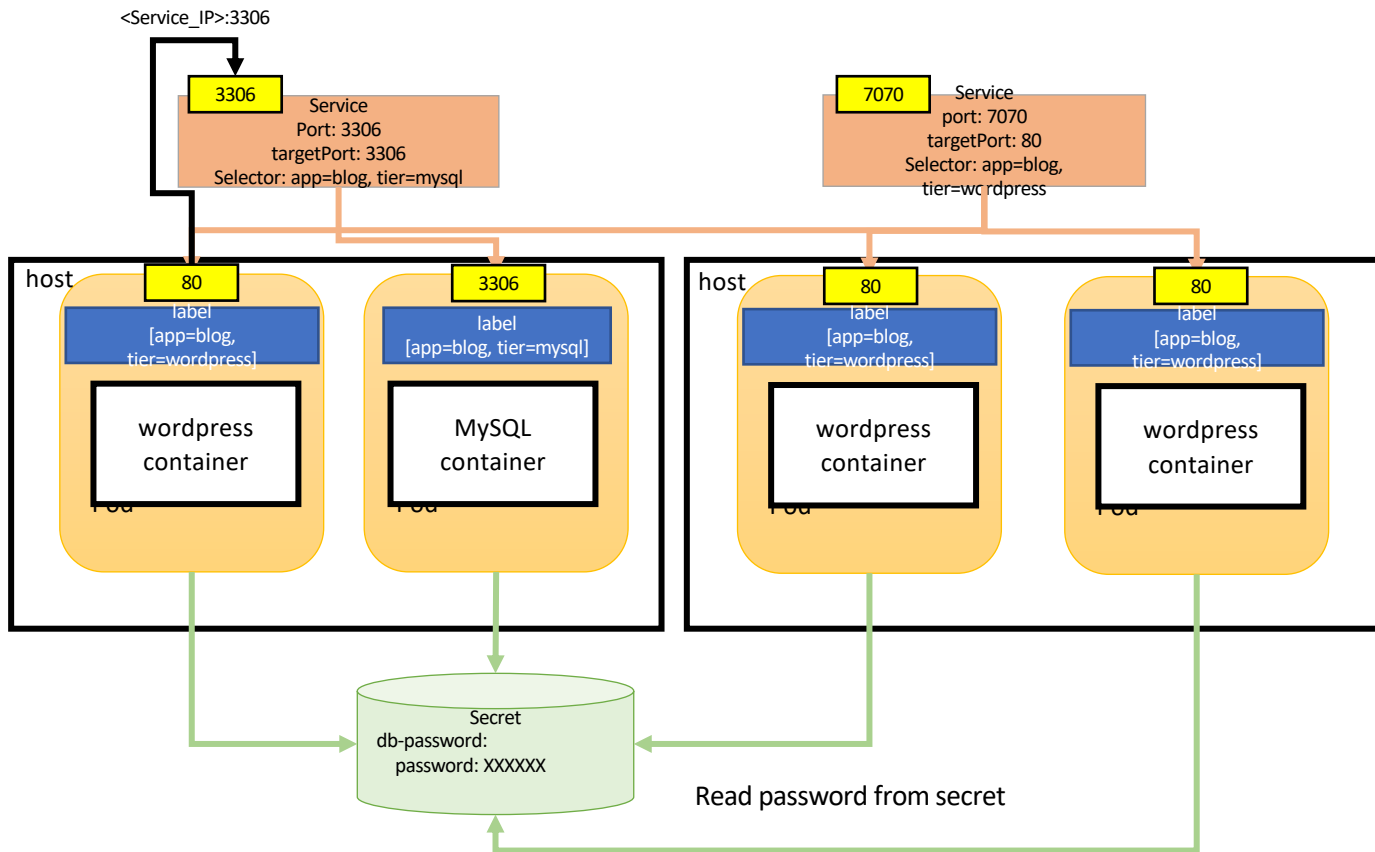
- 避免將機密資訊直接存放在Pod裡
- 透過mount / 環境變數 在 Pod裡使用
- 其實不安全
- 機密資料需先透過base64編碼
- 大小只能1MB

```
apiVersion: v1
kind: Secret
metadata:
  name: db-password
type: Opaque
data:
  password: UGFzc3dvcmQxMjM0
```

```
spec:
  containers:
    - name: nginx
      image: nginx:1.7.9
      volumeMounts:
        - name: db-secret
          mountPath: "/etc/db-secret"
          readOnly: true
      volumes:
        - name: db-secret
          secret:
            secretName: db-password
```



# Wordpress + MySQL的例子



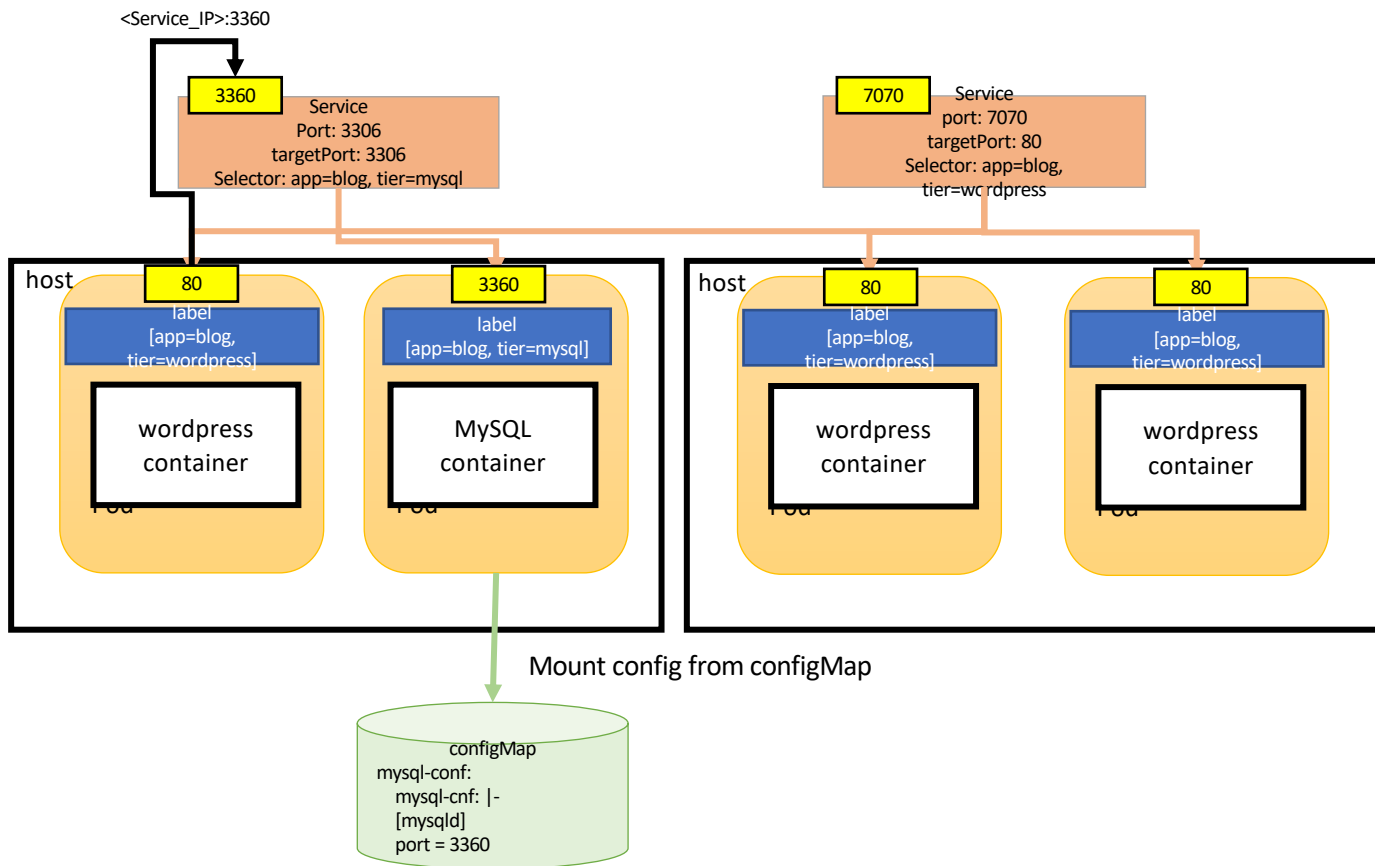
# ConfigMap

- 避免將非機密資訊直接存放在Pod裡
- 透過mount / 環境變數 在 Pod裡使用
- 其實不安全
- ~~• 機密資料需先透過base64編碼~~
- 大小只能1MB

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: mysql-conf
data:
  mysql-cnf: |-
    [mysqld]
    port = 3360
  mysql-port: "3360"
```

```
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          valueFrom:
            secretKeyRef:
              name: db-password
              key: password
      volumeMounts:
        - name: mysql-storage
          mountPath: /var/lib/mysql
        - name: mysql-conf
          mountPath: /etc/mysql/
  volumes:
    - name: mysql-storage
      hostPath:
        path: /tmp/data
    - name: mysql-conf
      configMap:
        name: mysql-conf
        items:
          - key: mysql-cnf
            path: my.cnf
```

# 在Wordpress + MySQL的例子中



# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap

核心元件

- 
- Ingress (請先安裝ingress-controller再進行範例)
  - Volume

開放第三方提供  
擴充元件

```
$ helm repo add bitnami https://charts.bitnami.com/bitnami
```

```
$ helm repo update
```

```
$ helm search repo bitnami/nginx
```

```
$ helm install ingress-controller --set kind=DaemonSet --set service.type=ClusterIP --set daemonset.useHostPort=true bitnami/nginx-ingress-controller
```

```
$ helm install ingress-controller --namespace ingress-controller --set nodeSelector."kubernetes\\.io/hostname"=master --set kind=Deployment --set service.type=ClusterIP --set hostNetwork=true --version 8.0.11 bitnami/nginx-ingress-controller
```

# Ingress

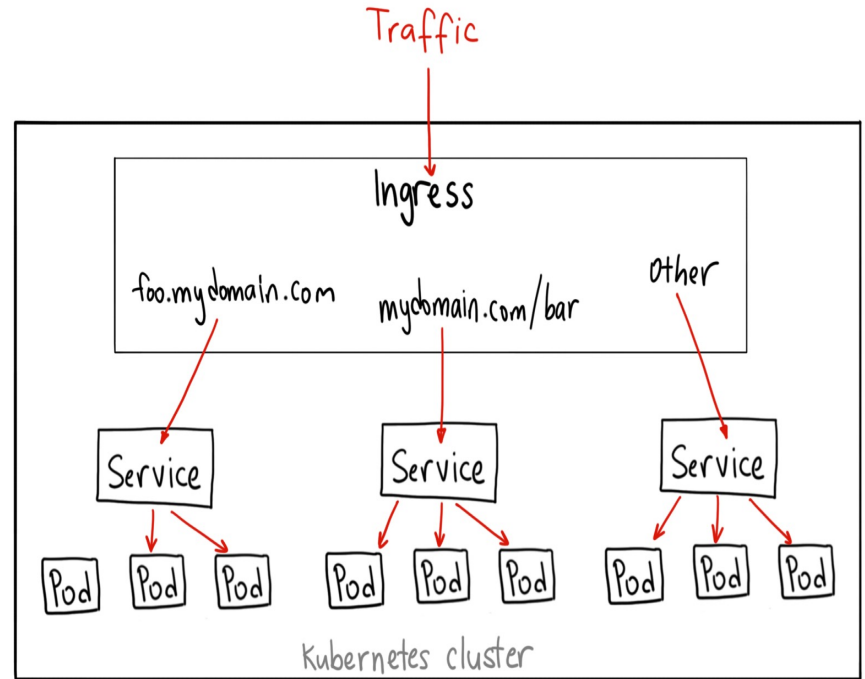
- 如何 expose service 給外部？
  - NodePort: 每個服務要對應一個nodePort
  - Load Balancer: 必需有cloud provider 支援
  - HostPort: 使用實體主機上的網路對外溝通
  - Ingress: 支援L7
- Ingress 包括
  - Ingress
    - 設定轉發的規則
  - Ingress-controller
    - 有不同的實作方法，nginx-controller是用nginx當proxy進行轉發
  - Service for ingress-controller
    - 讓外面存取不同服務的request都進入ingress-controller

```

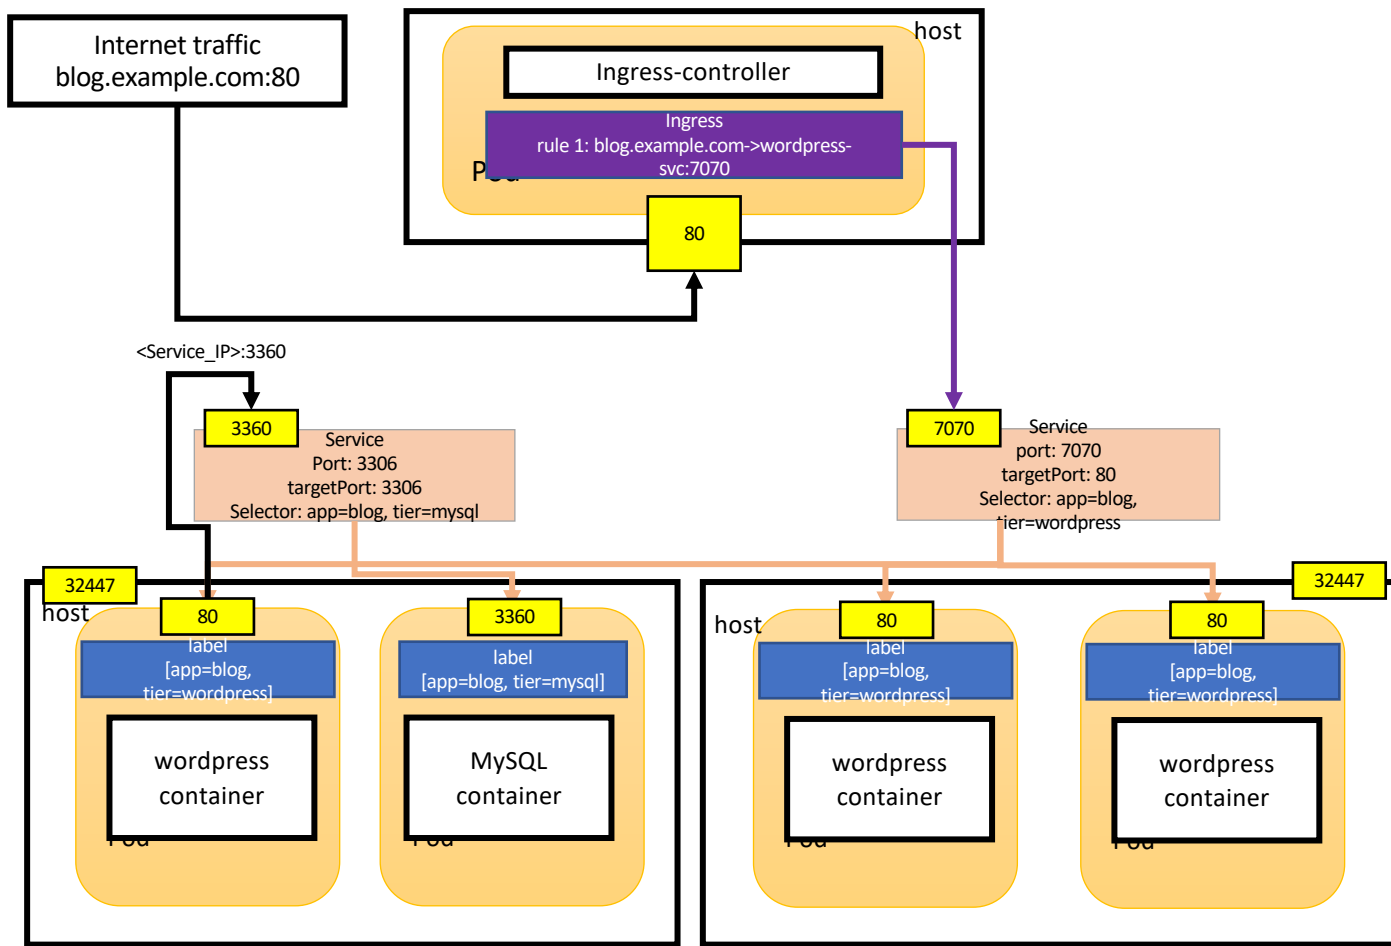
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  annotations:
    kubernetes.io/ingress.class: nginx
  name: wordpress-ingress
spec:
  rules:
  - host: blog.140-110-136-74.nip.io
    http:
      paths:
      - backend:
          service:
            name: wordpress-svc
            port:
              number: 7070
        pathType: ImplementationSpecific

```

Rule-1



# Ingress controller 採用hostPort



# Kubernetes Basic Concept

- Pod
- Label & Annotation
- Built-in workload
  - Deployment
- Service
- Secret & configMap
- Ingress
- Volume (請先安裝NFS server & nfs-provider再進行範例)



# Volume

- Volume & Persistence Volume
- Persistence Volume & Persistence Volume Claim
- Persistence Volume & StorageClass
- CSI Persistence Volume

# 為什麼需要Volume

- 容器裡的檔案是ephemeral，容器重啟後，資料就消失了
  - 要透過Volume進行持久化儲存
- Pod裡的container，可以共用volume
- Volume週期和pod一樣
- Kubernetes has different volume plugin

| Temp   | Local  | Network   |
|--|--|---|
| <ul style="list-style-type: none"><li>• emptyDir</li></ul> | <ul style="list-style-type: none"><li>• hostPath</li></ul> | <ul style="list-style-type: none"><li>• GlusterFS</li><li>• CephRBD</li><li>• gitRepo</li><li>• secret</li><li>• flocker</li><li>• gcePersistentDisk</li><li>• AWS ElasticBlockStore (EBS)</li><li>• NFS</li><li>• iSCSI</li><li>• Fibre Channel</li><li>• Cinder</li></ul> |

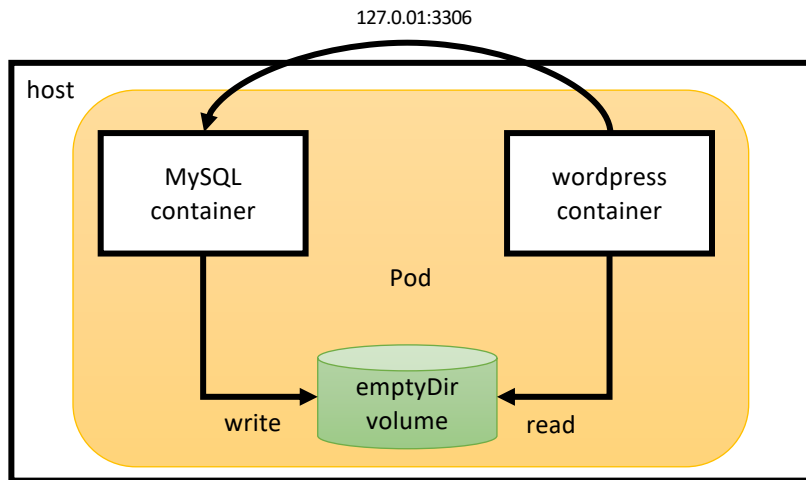
```

apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
  - name: mysql
    image: mysql:5.6
    env:
      - name: MYSQL_ROOT_PASSWORD
        value: Password1234
    volumeMounts:
      - name: mysql-storage
        mountPath: /var/lib/mysql
  - name: wordpress
    image: wordpress:4.8-apache
    env:
      - name: WORDPRESS_DB_HOST
        value: 127.0.0.1
      - name: WORDPRESS_DB_PASSWORD
        value: Password1234
    volumeMounts:
      - name: mysql-storage
        mountPath: /var/lib/wordpress
  volumes:
    - name: mysql-storage
      emptyDir: {}

```

## • 使用emptyDir

- 生命週期和pod相同
  - 當pod被移除時，emptyDir也會被清空
  - 但Container 重啟，資料仍會保留
- 同一個pod的container間，可以共享資料



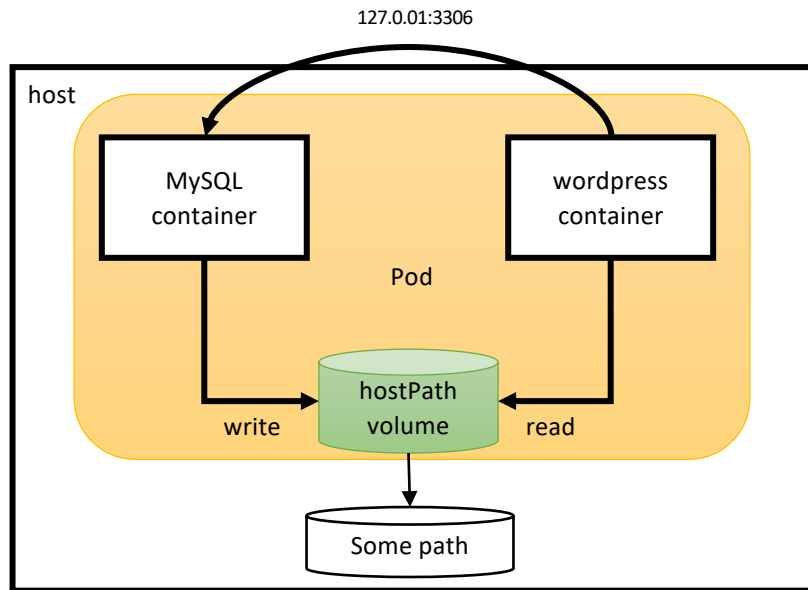
```

apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
  - name: mysql
    image: mysql:5.6
    env:
    - name: MYSQL_ROOT_PASSWORD
      value: Password1234
    volumeMounts:
    - name: mysql-storage
      mountPath: /var/lib/mysql
  - name: wordpress
    image: wordpress:4.8-apache
    env:
    - name: WORDPRESS_DB_HOST
      value: 127.0.0.1
    - name: WORDPRESS_DB_PASSWORD
      value: Password1234
    volumeMounts:
    - name: mysql-storage
      mountPath: /var/lib/wordpress
  volumes:
  - name: mysql-storage
    hostPath:
      path: /tmp/data

```

## • 使用hostPath

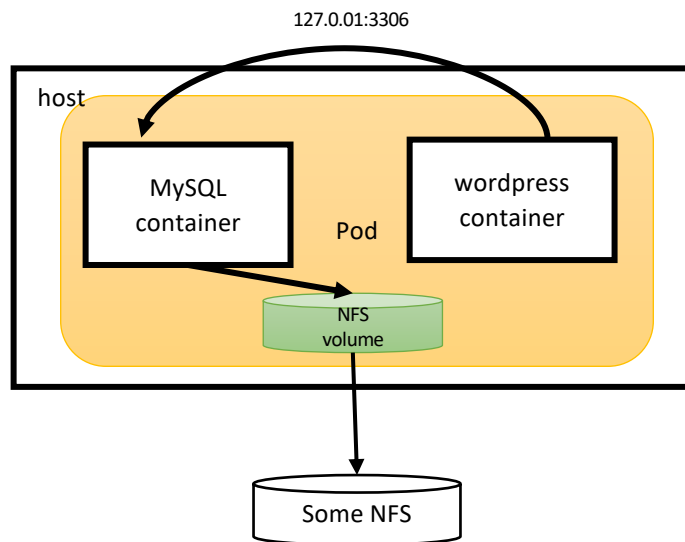
- 將host上的某個目錄掛載至Pod
- 資料會保存在host上，重啟pod不會清空資料
- 但container 移動到別的host上，資料遺失



# Wordpress + MySQL的例子

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-storage
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
  volumes:
    - name: mysql-volumes
      nfs:
        server: 192.168.2.31
        path: /nfs-data
```

- 使用nfs



使用者

3. 在YAML裡使用

2. 設定好告知使用者

管理者

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-storage
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
  volumes:
    - name: mysql-volumes
      nfs:
        server: 192.168.2.31
        path: /nfs-data
```

1. 讓K8S支援

儲存設備業者

# Persistence Volume (PV) & Persistence Volume Claim(PVC)

- Abstraction of Physical storage
  - Separate **Dev**elopment & **Op**erator
- PersistentVolume (PV)
  - A networked storage provisioned **by an administrator**
  - Operator view
- PersistentVolumeClaim (PVC)
  - Request for storage **from a user**
  - Development view

```


apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data

```

Administrator



Registers PVs in the pool




Administrator owned

Developer



Claims a PV from the pool

References claim in pod



Developer owned

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Gi

```

```

apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-volumes
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
      volumes:
        - name: mysql-volumes
          persistentVolumeClaim:
            claimName: nfs-pvc

```



使用者

3-0. 宣告想用的儲存  
3-1. 綁定成功後使用



```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Gi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-volumes
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
  volumes:
    - name: mysql-volumes
      persistentVolumeClaim:
        claimName: nfs-pvc
```

管理者

2. 建立可用的  
Volume



```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data
```

1. 讓K8S支援



儲存設  
備業者

# Storage Class(SC)

- Before User request a PVC, administrator need to create PV in advanced.
  - Static Provisioning
- If storage volumes can be created on-demand
  - Dynamic Provisioning
  - use storage class to enable Dynamic Provisioning
  - Provisioner is required for provisioning PVs
  - Specify SC in PVC
    - If SC is not specified in PVC, default SC is used
    - If there is no match SC, PVC is pending

```


apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data

```

Administrator



Registers PVs in the pool




Administrator owned

Developer



Claims a PV from the pool

References claim in pod




Developer owned

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Gi

```

Create PVs on-demand



```

apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-volumes
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
      volumes:
        - name: mysql-volumes
          persistentVolumeClaim:
            claimName: nfs-pvc

```

使用者

- 3-0. 宣告想用的儲存  
3-1. 透過Provisioner建立volume  
3-2. 綁定成功後使用

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Gi
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-volumes
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
      volumes:
        - name: mysql-volumes
          persistentVolumeClaim:
            claimName: nfs-pvc
```

管理者

2. 安裝provisioner

Storage Class

provisioner

- 1-0. 讓K8S支援  
1-1. 開發Provisioner

儲存設備業者

# CSI Persistence Volume

- 透過CSI, 各家廠商可以實作自己的Storage Driver

NFS 的PV 與CSI版本的NFS PV用法

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-nfs
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteMany
  mountOptions:
    - hard
    - nfsvers=4.1
  csi:
    driver: nfs.csi.k8s.io
    volumeHandle: unique-volumeid
    volumeAttributes:
      server: 192.168.2.201
      share: /
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: new-pv
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 5Gi
  csi:
    driver: ch.ctrox.csi.s3-driver
    volumeAttributes:
      secretNamespace: default
      secretName: user-s3-secret
    volumeHandle: ogre0403
```

CSI版本的NFS 與 S3 PV用法

使用者

- 3-0. 宣告想用的儲存  
3-1. 透過Provisioner建立volume  
3-2. 綁定成功後使用

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Gi
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.2.31
    path: /nfs-data
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-wordpress-pod
spec:
  containers:
    - name: mysql
      image: mysql:5.6
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: Password1234
      volumeMounts:
        - name: mysql-volumes
          mountPath: /var/lib/mysql
    - name: wordpress
      image: wordpress:4.8-apache
      env:
        - name: WORDPRESS_DB_HOST
          value: 127.0.0.1
        - name: WORDPRESS_DB_PASSWORD
          value: Password1234
  volumes:
    - name: mysql-volumes
      persistentVolumeClaim:
        claimName: nfs-pvc
```

管理者

2. 安裝provisioner

Storage Class

provisioner

儲存設備業者

- 1-0. 開發csi driver, 不需整合至K8S  
1-1. 開發Provisioner

# Summary

- Pod透過volume持久化資料
- 透過PVC，提供volume的抽象概念，PVC需綁定至底層的一個PV
  - 使用者不需要了解Volume是從那裡來的
  - PV需事先由管理者建立
- 透過storage class與Provisioner, 可以自動產生PV
- 透過CSI, 設備廠商可以自行擴充讓K8S支援新的產品

# Exercise

- 撰寫yaml檔，部署nginx
  - 建立一個deployment.yaml，使用nginx image
  - svc.yaml, 採用nodeport
  - 可以由本機連到NGINX的deployment